

# Reliable Data Transfer Protocol

CS 3516

B-term 2018

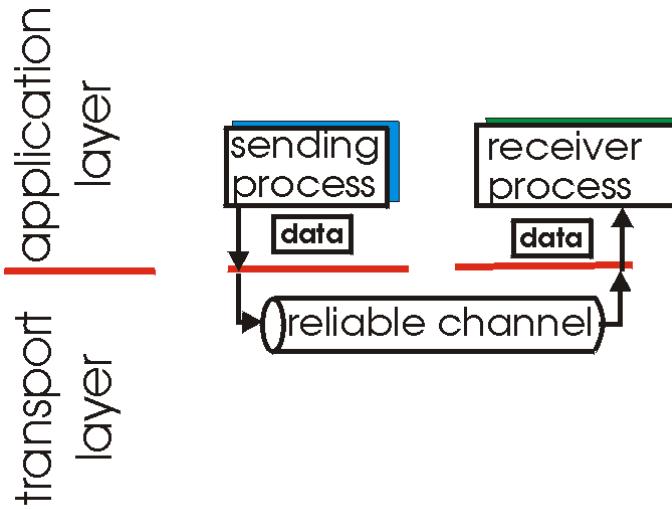
Instructor: Krishna Venkatasubramanian

# Announcements

- Project 1 due on Friday
- Project 2 and Lab2 will be out on Friday
- Quiz 2 on Friday
  - Will include questions from topics in chapter 1 (metrics), application layer, along with topics from today

# Principles of reliable data transfer

- important in application, transport, link layers
  - top-10 list of important networking topics!

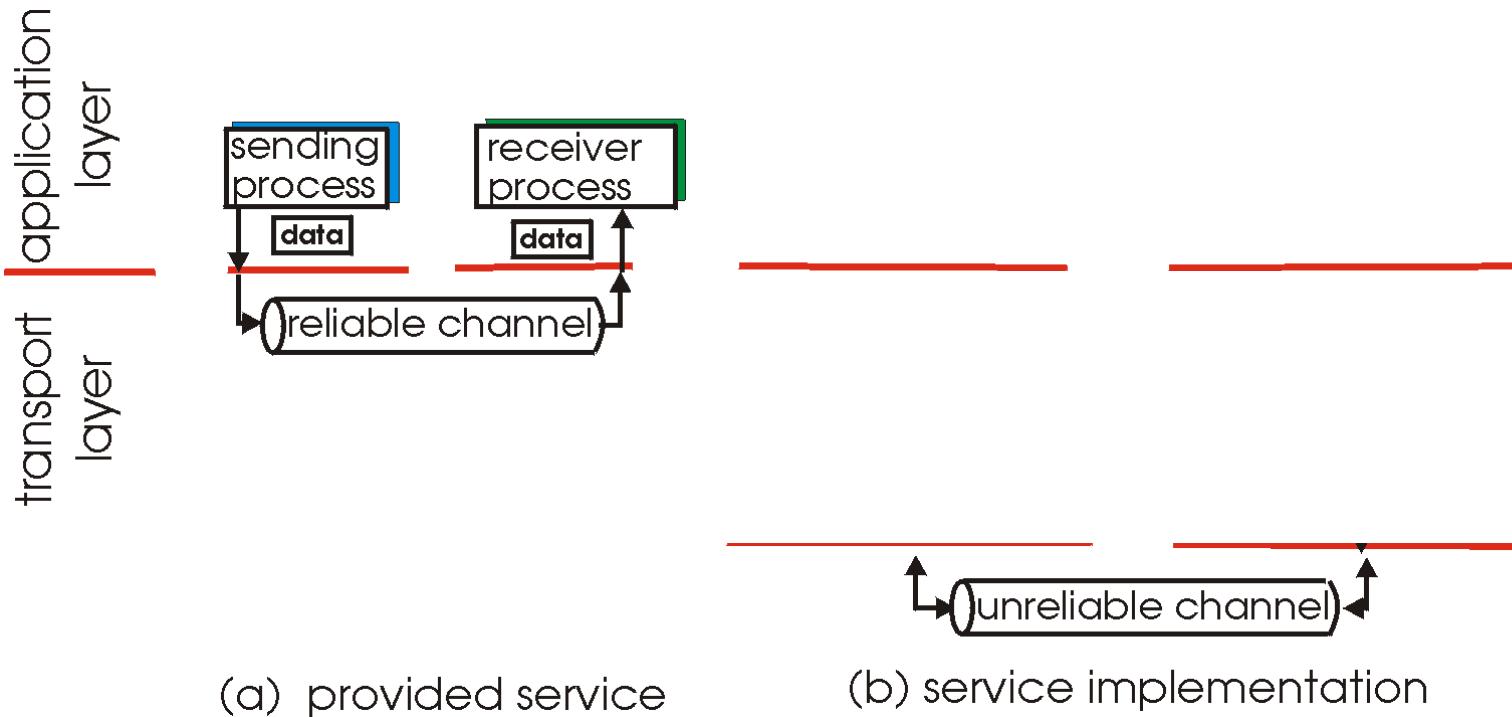


(a) provided service

- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of reliable data transfer

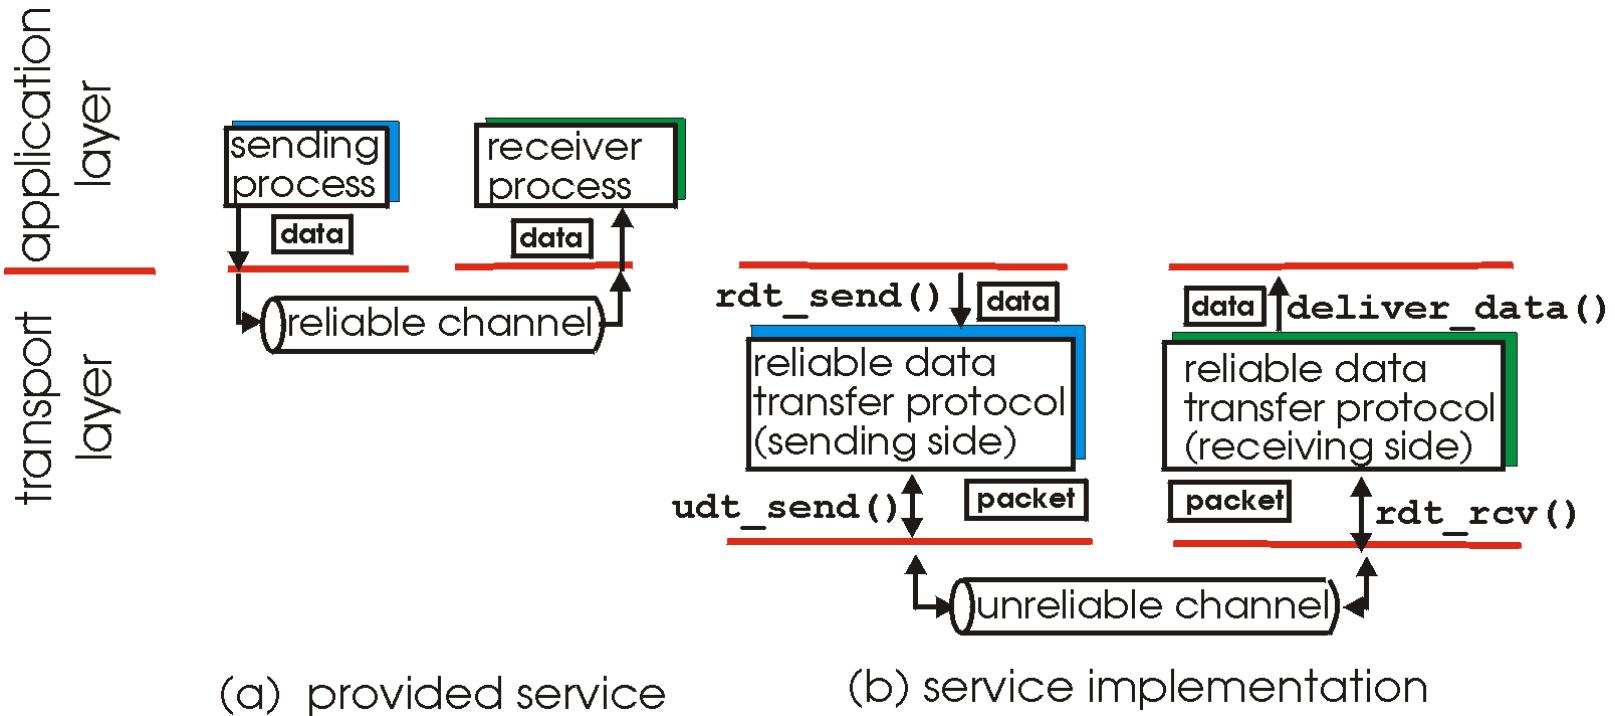
- important in application, transport, link layers
  - top-10 list of important networking topics!



- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of reliable data transfer

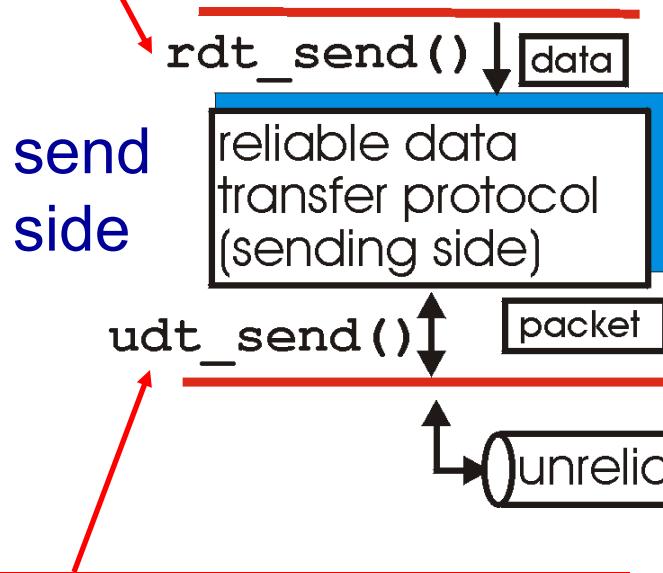
- important in application, transport, link layers
  - top-10 list of important networking topics!



- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

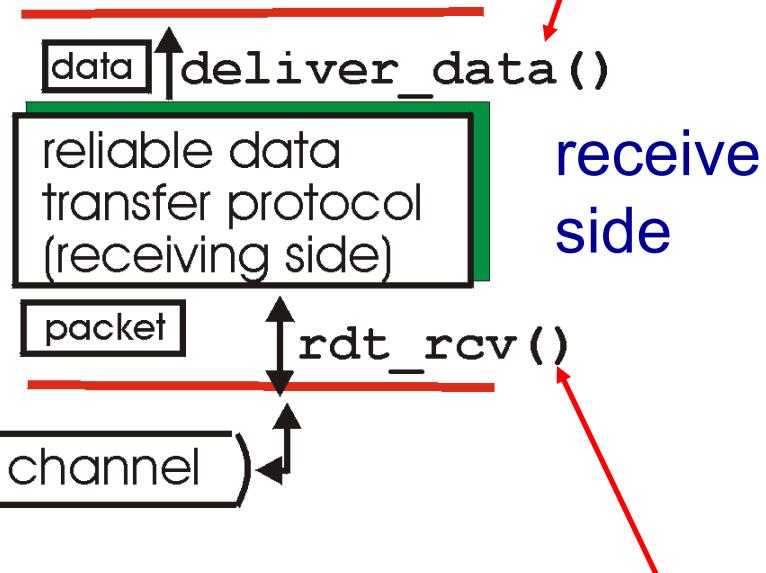
# Reliable data transfer: getting started

**rdt\_send()** : called from above, (e.g., by app.). Passed data to deliver to receiver upper layer



**udt\_send()** : called by rdt, to transfer packet over unreliable channel to receiver

**deliver\_data()** : called by **rdt** to deliver data to upper



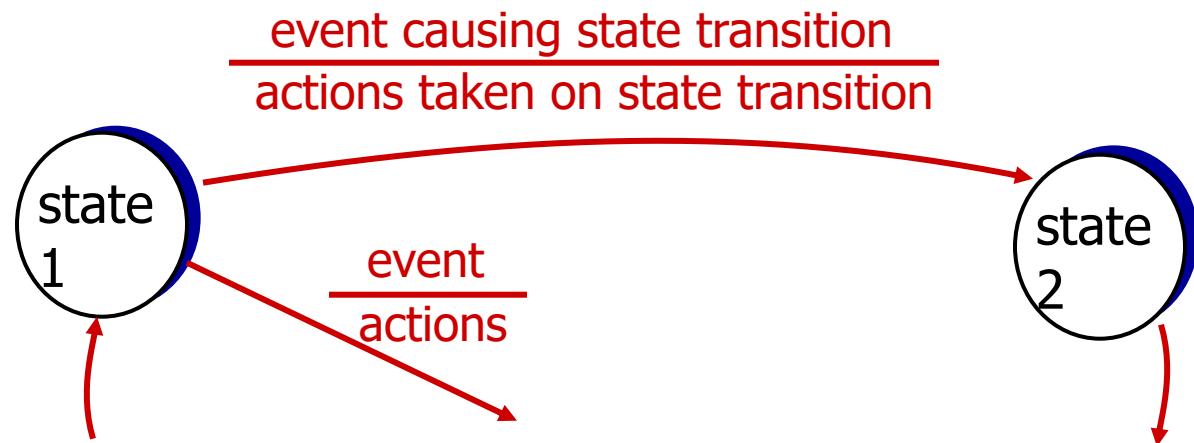
**rdt\_rcv()** : called when packet arrives on rcv-side of channel

# Reliable data transfer: getting started

we'll:

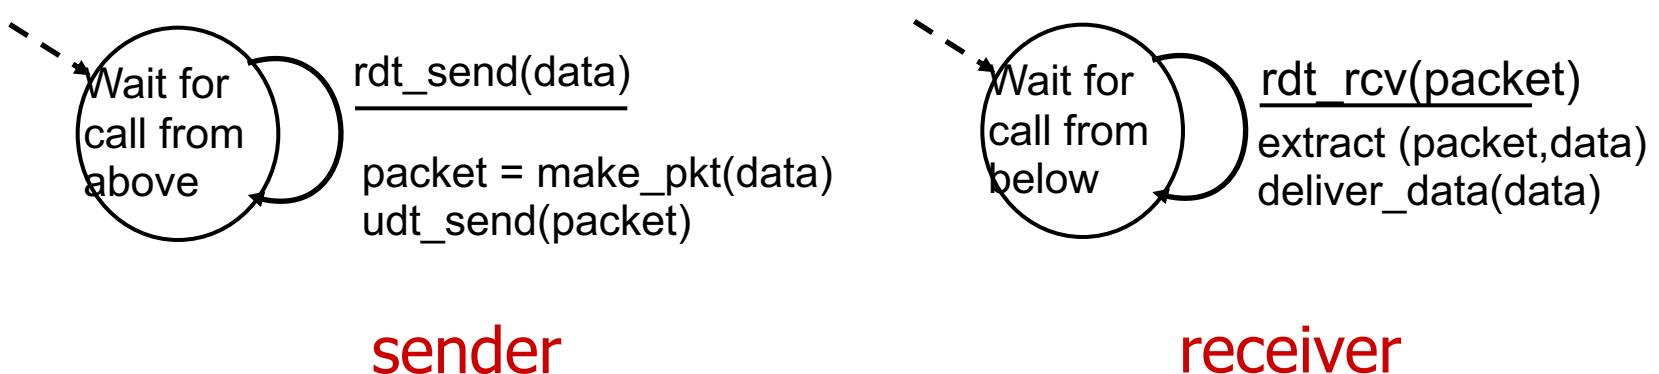
- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
  - but control info will flow on both directions!
- THREE** main issues in reliable data transfer:
  - Packet Corruption
  - Packet Loss
  - Packet reordering (NOT CONSIDERED)
- use finite state machines (FSM) to specify sender, receiver

**state:** when in this “state” next state uniquely determined by next event



# rdt 1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel



## rdt2.0: channel with bit errors (still no packet loss)

- underlying channel may flip bits in packet
  - checksum to detect bit errors
- *the question: how to recover from errors:*

*How do humans recover from “errors”  
during conversation?*

# rdt2.0: channel with bit errors

- underlying channel may flip bits in packet
  - checksum to detect bit errors
- *the question: how to recover from errors:*
  - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
  - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK
- new mechanisms in rdt2.0 (beyond rdt1.0):
  - error detection
  - feedback: control msgs (ACK,NAK) from receiver to sender for every packet
  - retransmission of messages by sender for which NAK is received
- Called **Automatic Repeat Request (ARQ)** protocols

# rdt2.0: FSM specification

rdt\_send(data)

sndpkt = make\_pkt(data, checksum)

udt\_send(sndpkt)

receiver

Wait for  
call from  
above

Wait for  
ACK or  
NAK

rdt\_rcv(rcvpkt) &&  
isNAK(rcvpkt)

udt\_send(sndpkt)

rdt\_rcv(rcvpkt) && isACK(rcvpkt)

$\Lambda$

sender

rdt\_rcv(rcvpkt) &&  
corrupt(rcvpkt)

udt\_send(NAK)

Wait for  
call from  
below

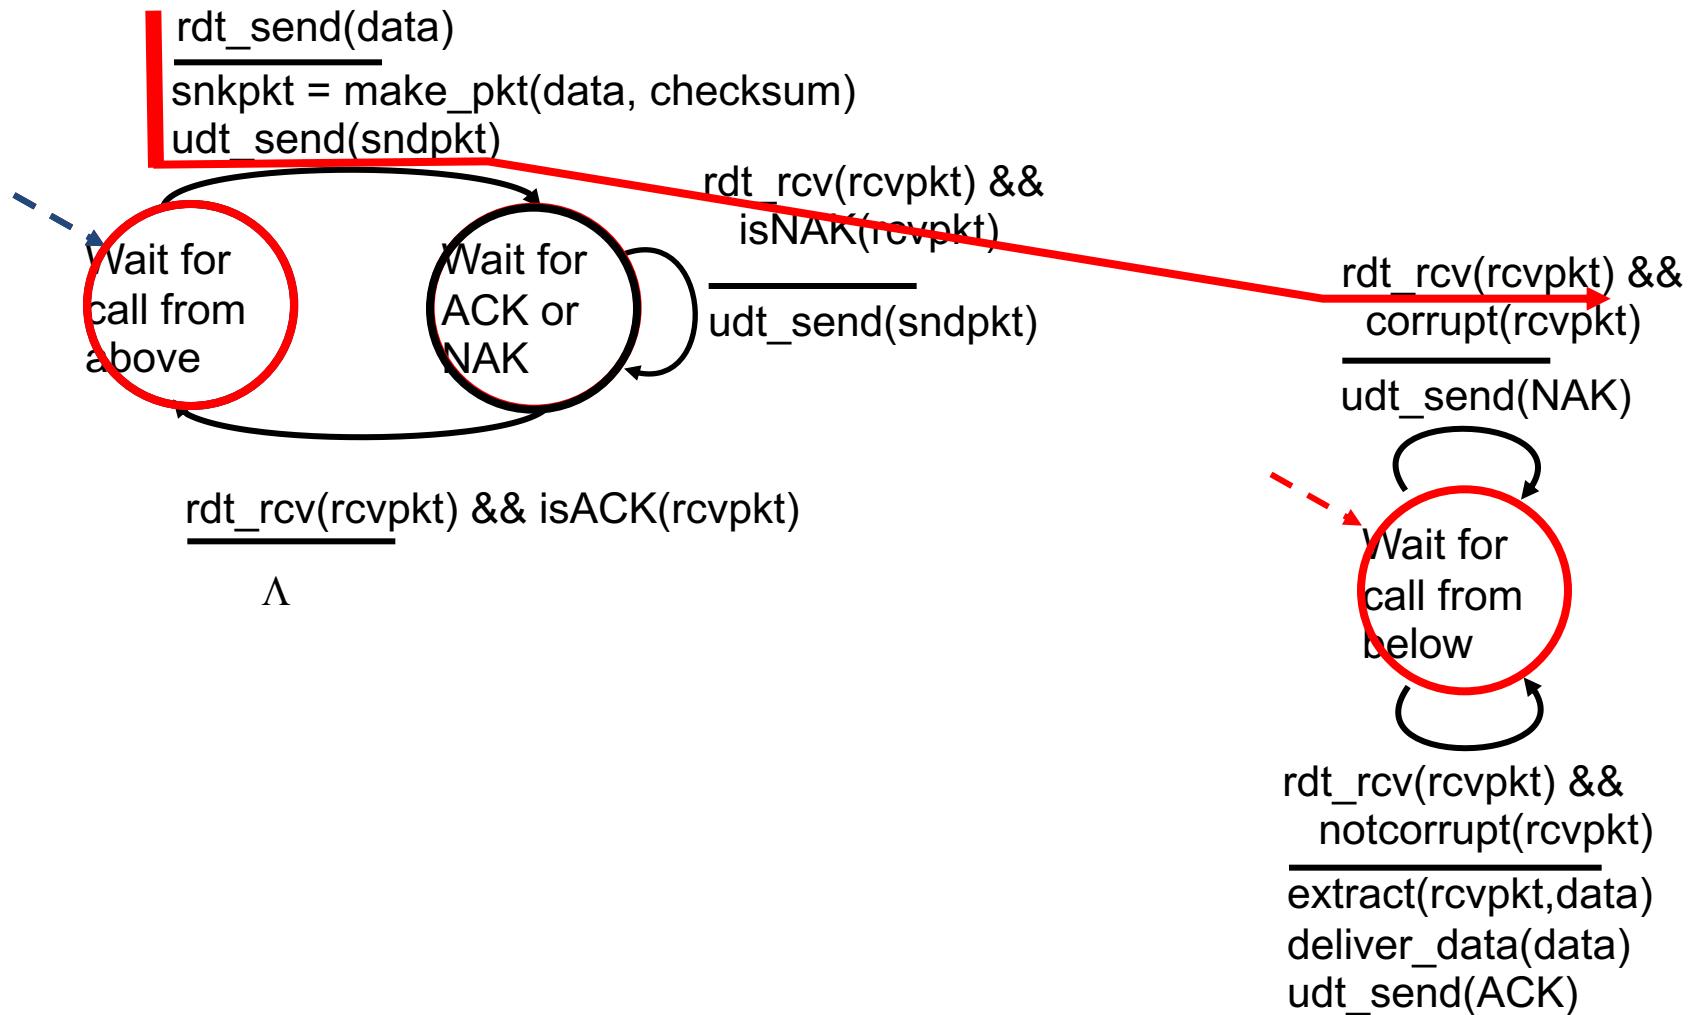
rdt\_rcv(rcvpkt) &&  
notcorrupt(rcvpkt)

extract(rcvpkt,data)

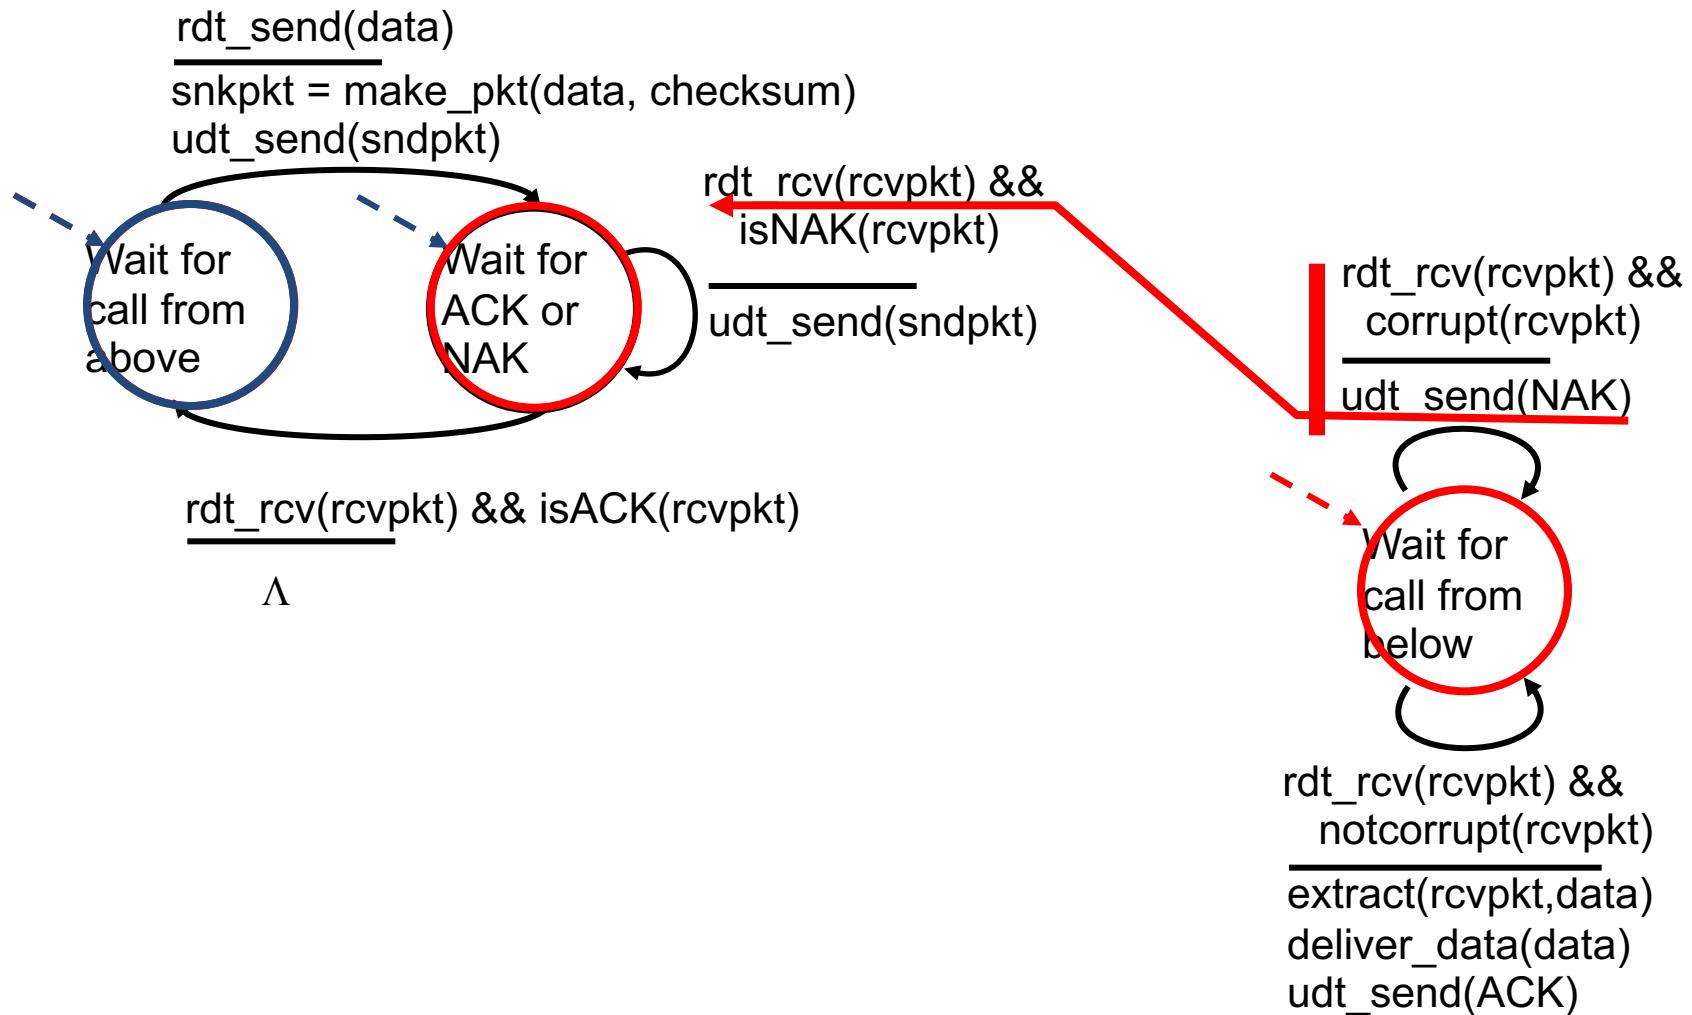
deliver\_data(data)

udt\_send(ACK)

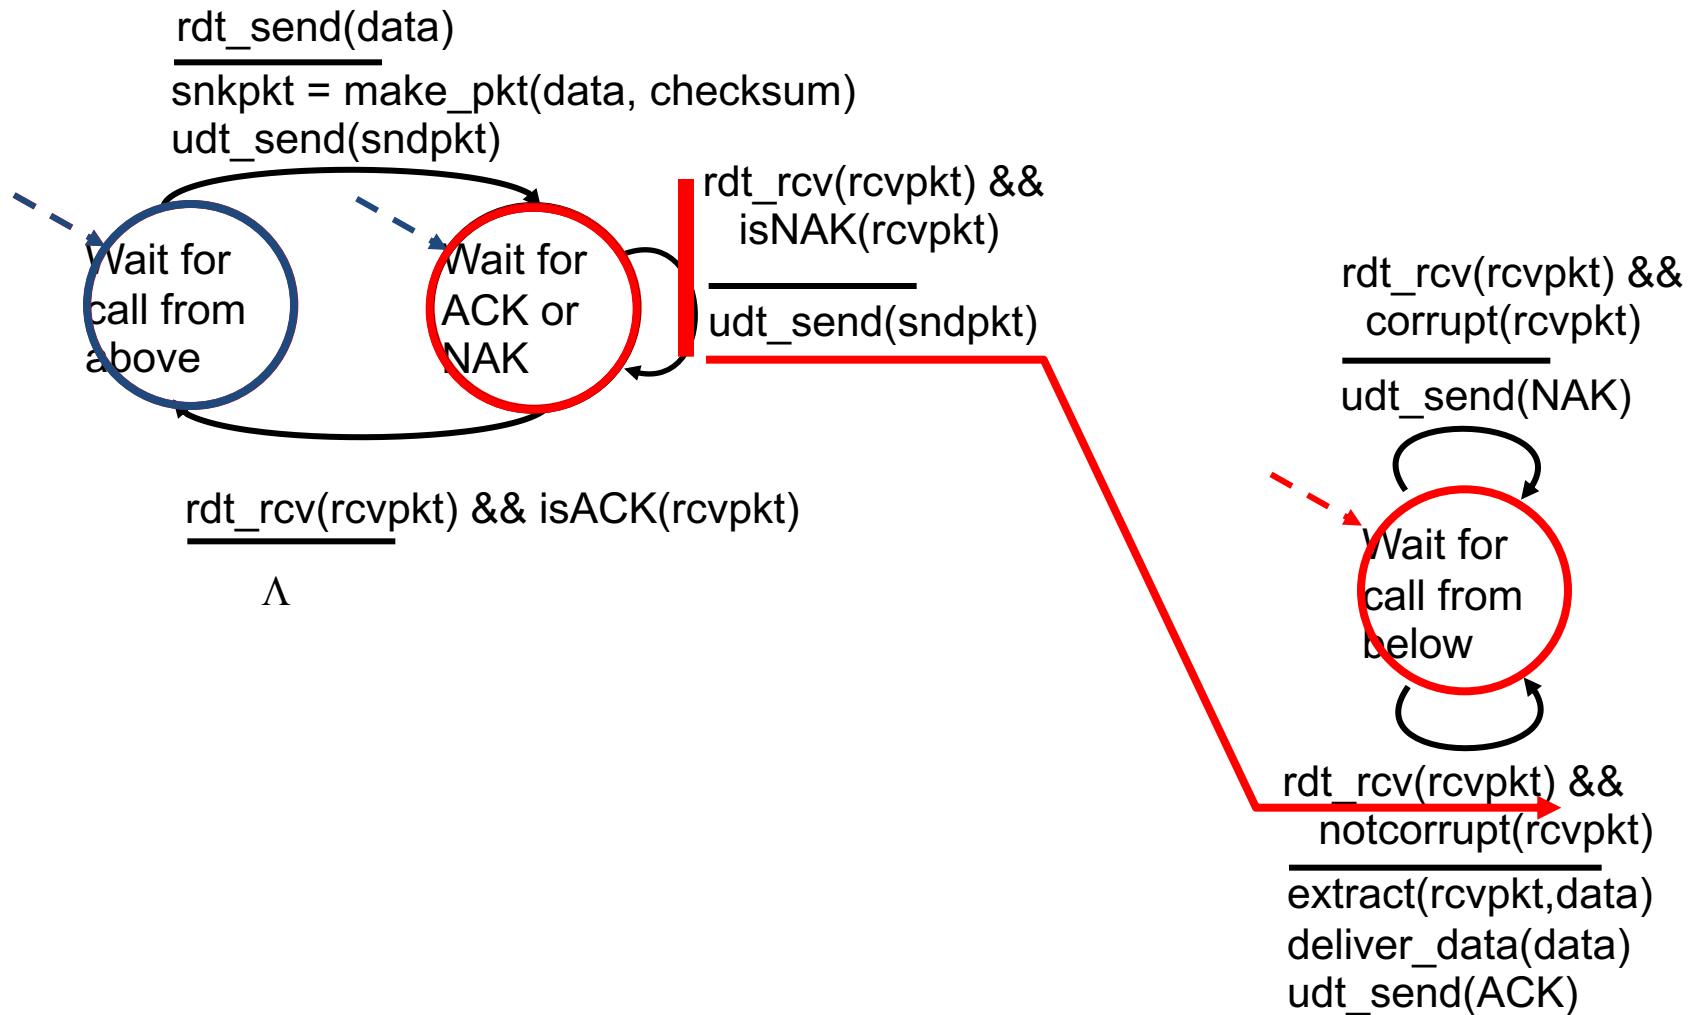
# rdt2.0: error scenario



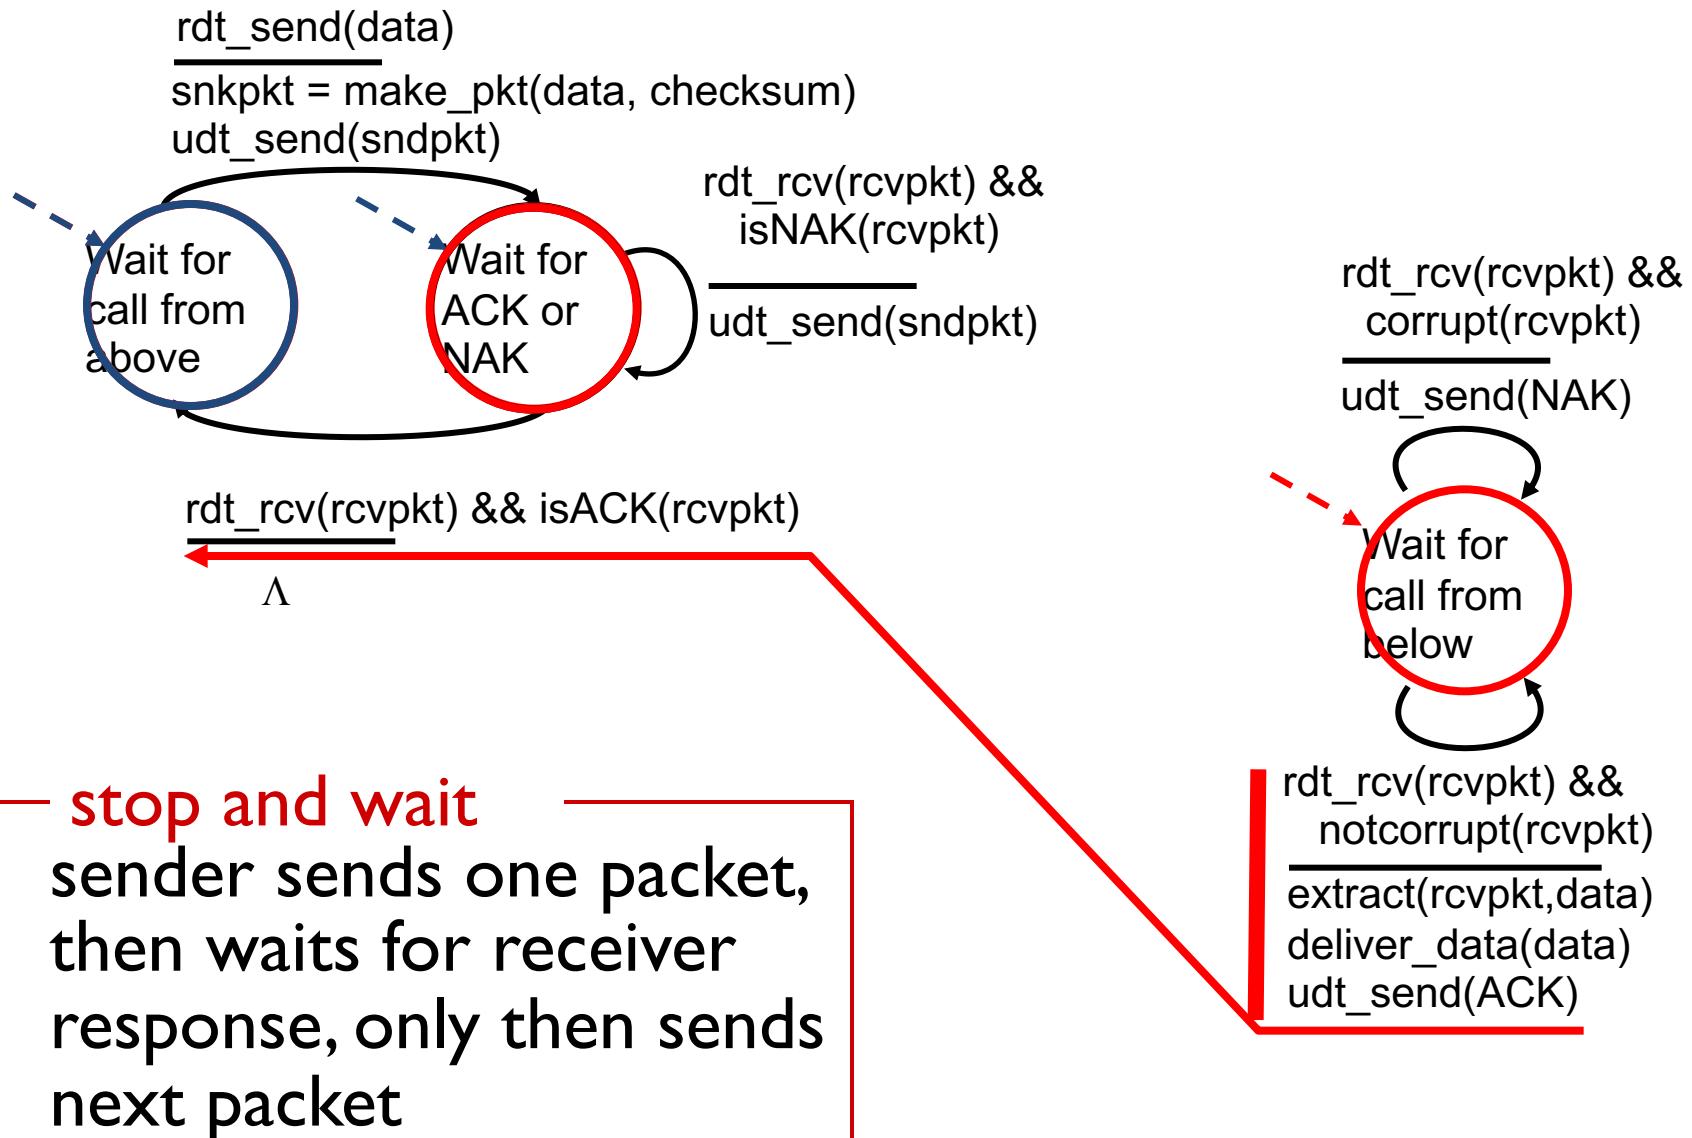
# rdt2.0: error scenario



# rdt2.0: error scenario



# rdt2.0: error scenario



# rdt2.0 has a fatal flaw!

what happens if  
ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit:  
**possible duplicate**

# rdt2.0 has a fatal flaw!

## what happens if ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit:  
**possible duplicate**

## handling duplicates:

- sender retransmits current pkt if ACK/NAK corrupted
- sender adds *sequence number* to each pkt
- receiver discards (doesn't deliver up) duplicate pkt

# rdt2.1: dealing with ack/nak errors

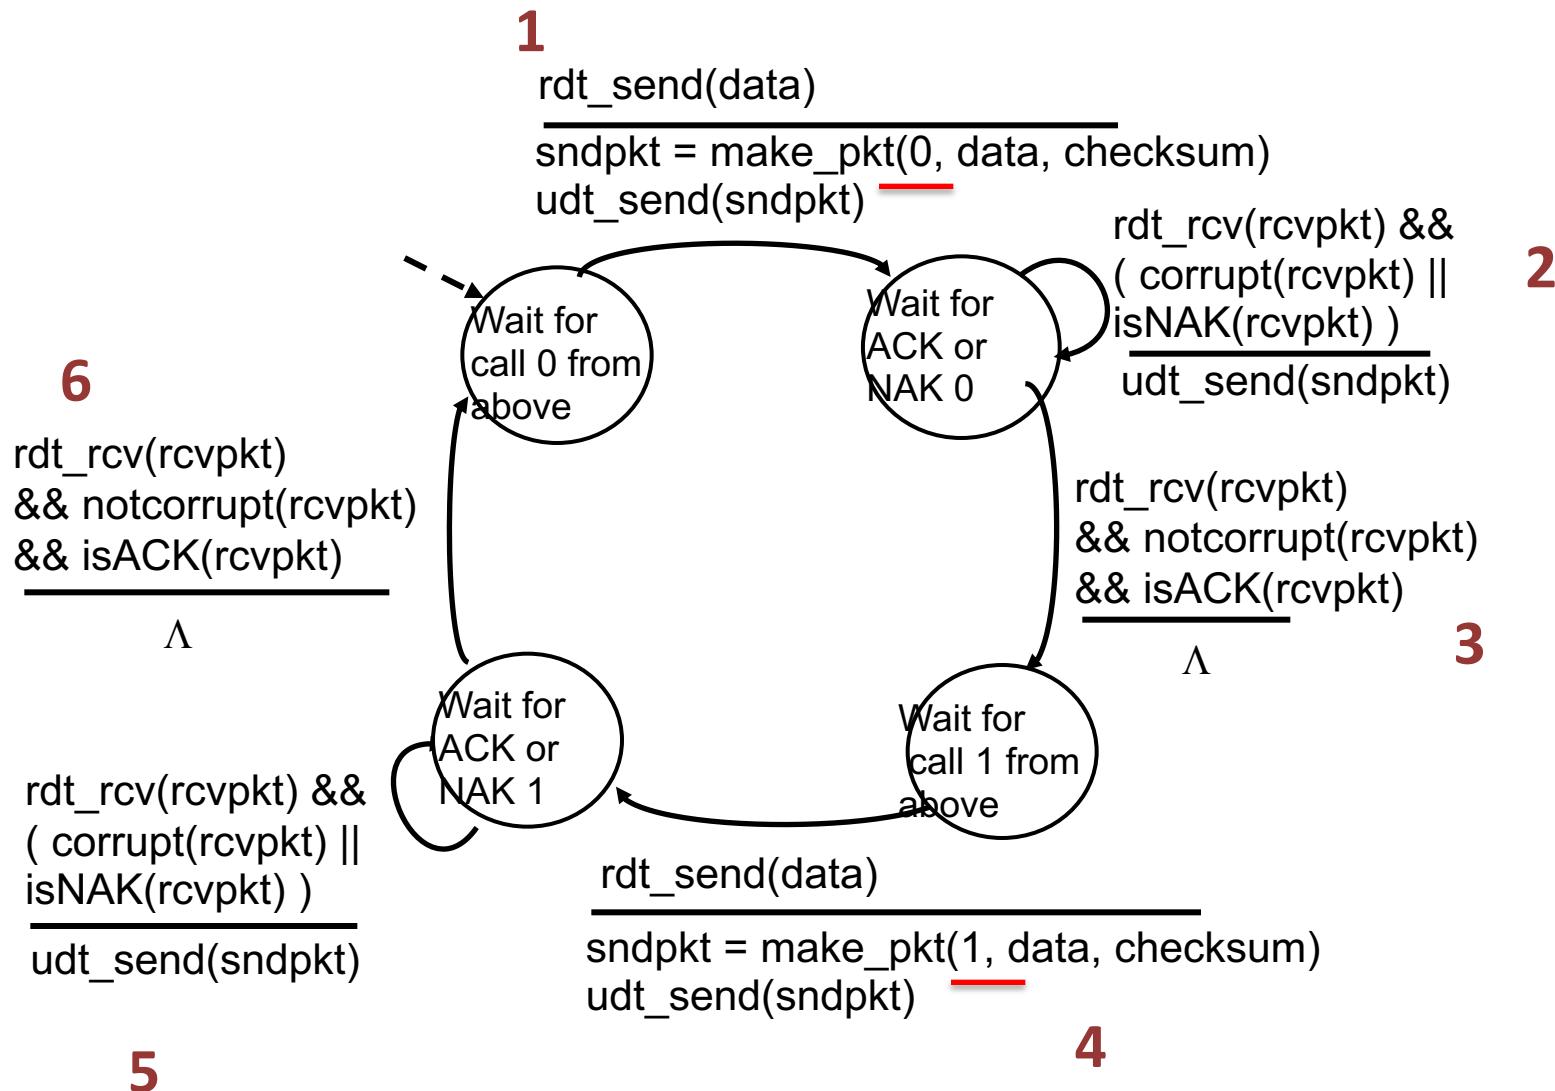
## sender:

- seq # added to pkt
- two seq. #'s (0,1) will suffice. **Why?**
- must check if received ACK/NAK corrupted
- twice as many states
  - state must “remember” whether “expected” pkt should have seq # of 0 or 1

## receiver:

- must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver cannot know if its last ACK/NAK received OK at sender
  - Only way to know that is if the same packet # was received more than once

# rdt2.1: sender, handles garbled ACK/NAKs



# rdt2.1: receiver, handles garbled ACK/NAKs

3

rdt\_rcv(rcvpkt) && (corrupt(rcvpkt)

  sndpkt = make\_pkt(NAK, chksum)  
  udt\_send(sndpkt)

rdt\_rcv(rcvpkt) &&  
  not corrupt(rcvpkt) &&  
  has\_seq1(rcvpkt)

  sndpkt = make\_pkt(ACK, chksum)  
  udt\_send(sndpkt)

1

rdt\_rcv(rcvpkt) && notcorrupt(rcvpkt)  
  && has\_seq0(rcvpkt)

---

  extract(rcvpkt,data)  
  deliver\_data(data)  
  sndpkt = make\_pkt(ACK, chksum)  
  udt\_send(sndpkt)

4

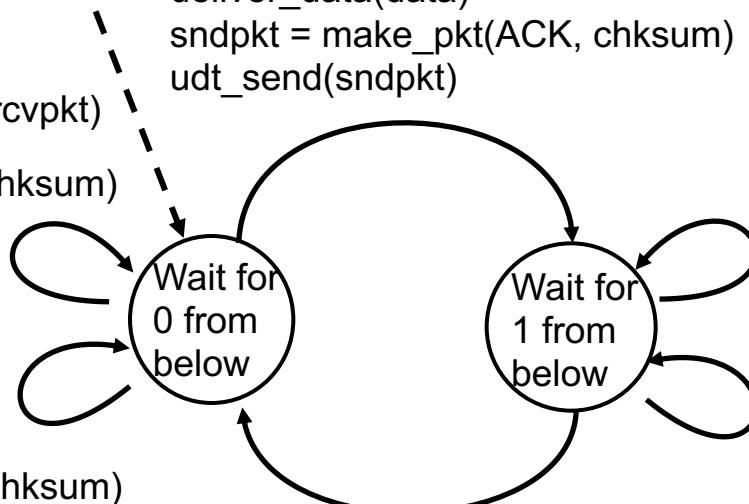
rdt\_rcv(rcvpkt) && (corrupt(rcvpkt)

  sndpkt = make\_pkt(NAK, chksum)  
  udt\_send(sndpkt)

rdt\_rcv(rcvpkt) &&  
  not corrupt(rcvpkt) &&  
  has\_seq0(rcvpkt)

  sndpkt = make\_pkt(ACK, chksum)  
  udt\_send(sndpkt)

2

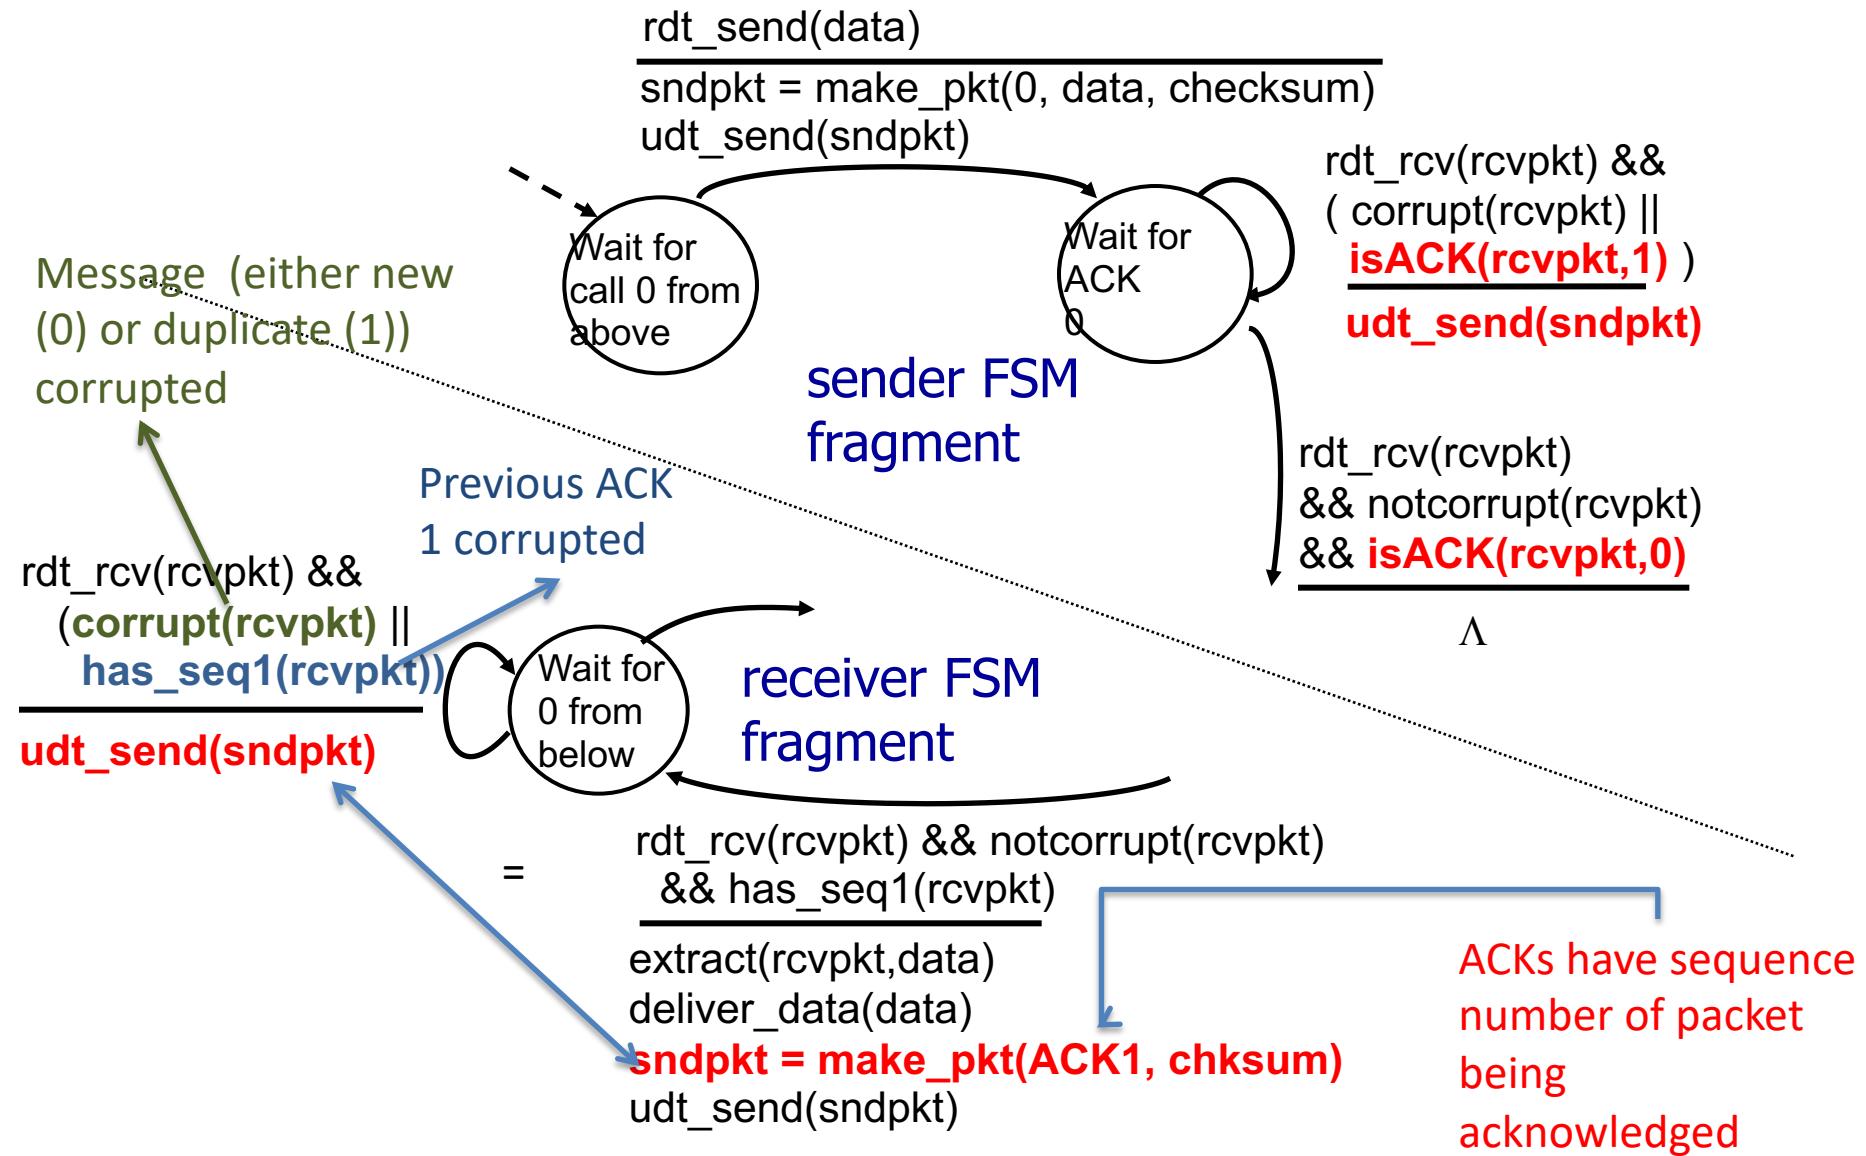


What does  
this mean?

## rdt2.2: a NAK-free (ACK-only) protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must *explicitly* include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

# rdt2.2: sender, receiver fragments



# rdt3.0: channels with errors and loss/delay

## new assumption:

underlying channel can also lose packets (data, ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

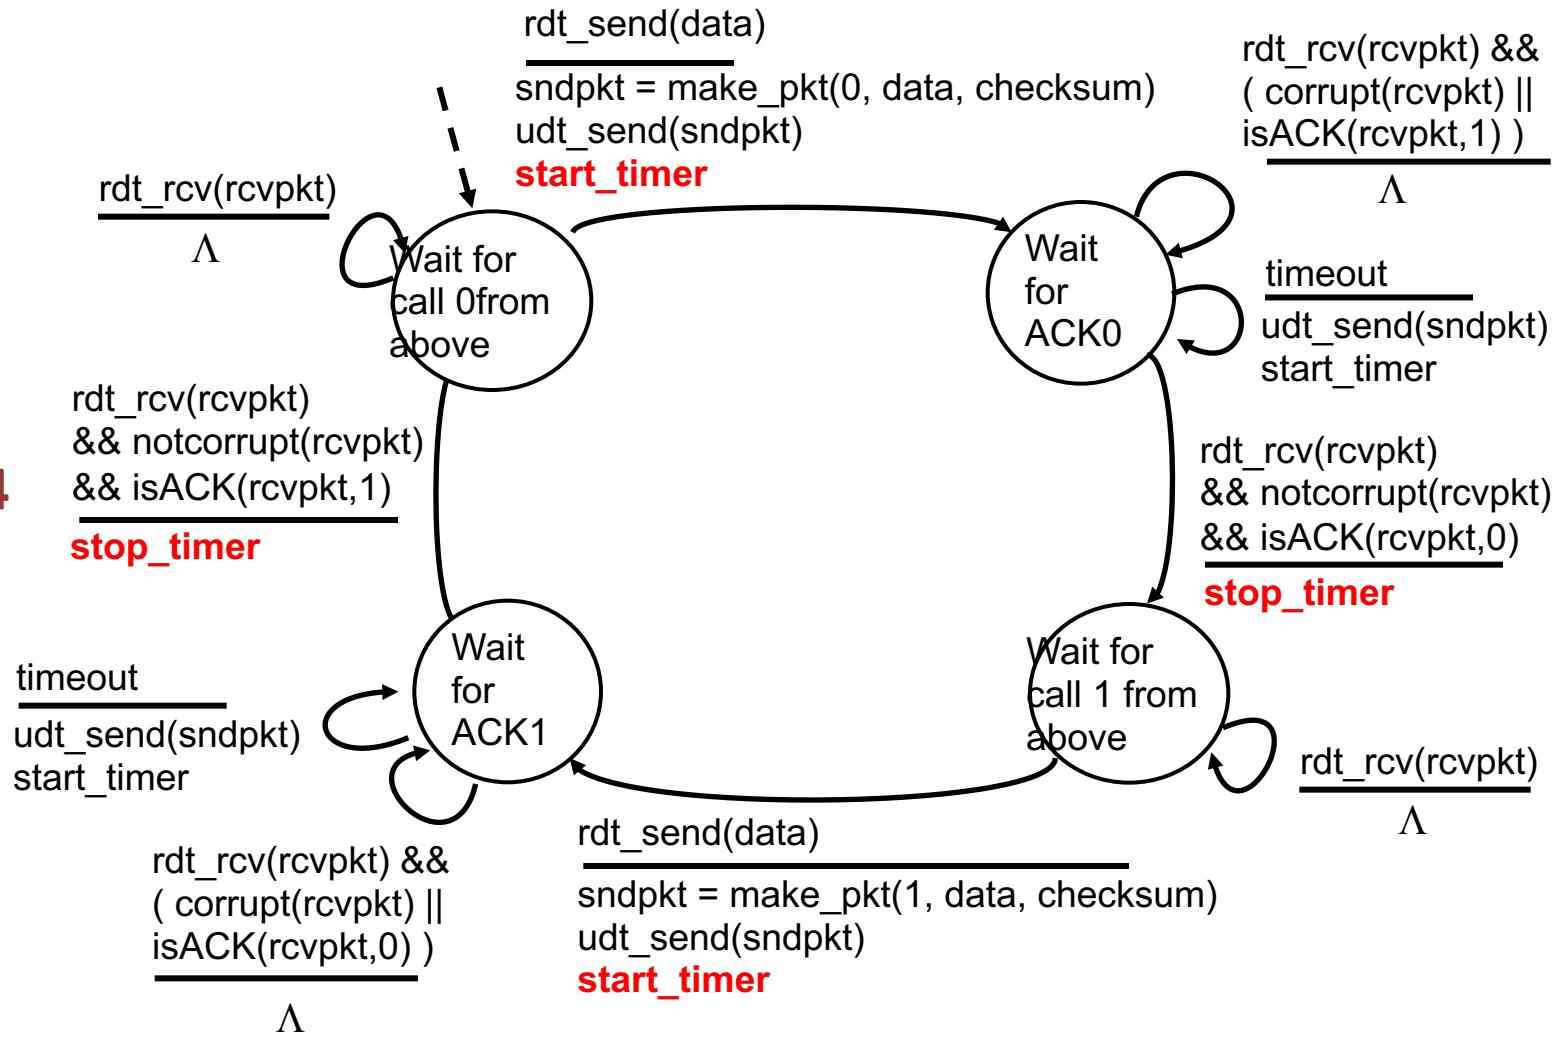
## approach: sender waits

“reasonable” amount of time for ACK

- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but seq. #’s already handles this
  - receiver must specify seq # of pkt being ACKed
- requires countdown timer
  - **For each PACKET**

# rdt3.0 sender

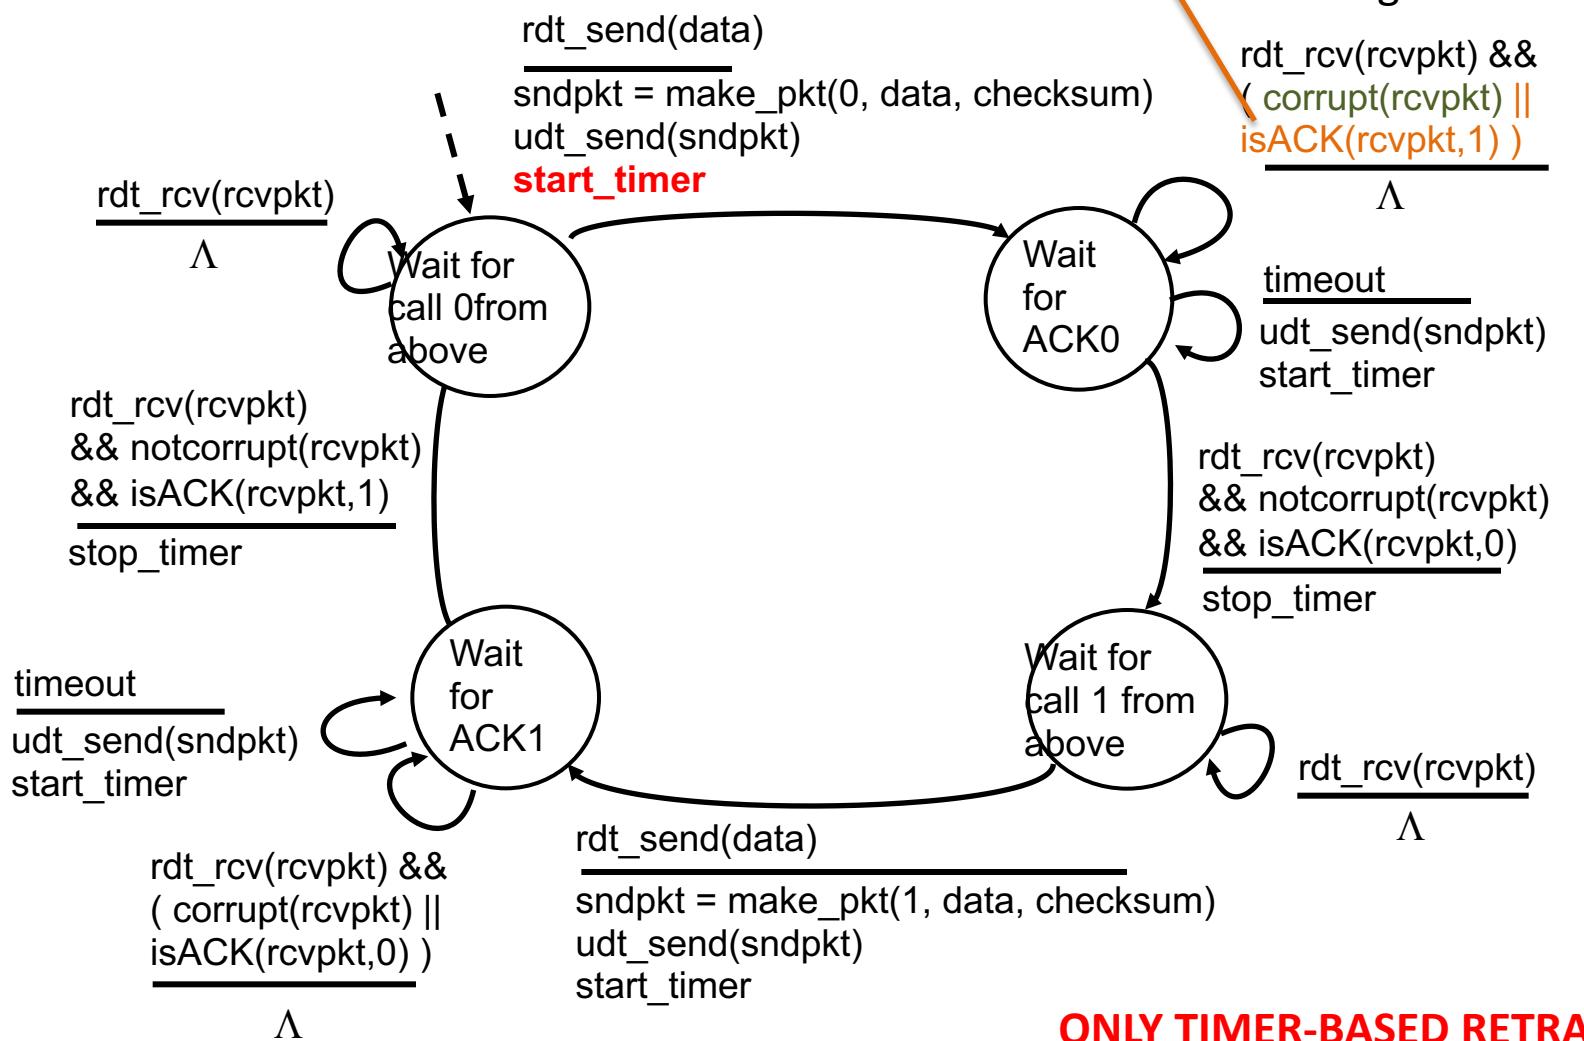
1



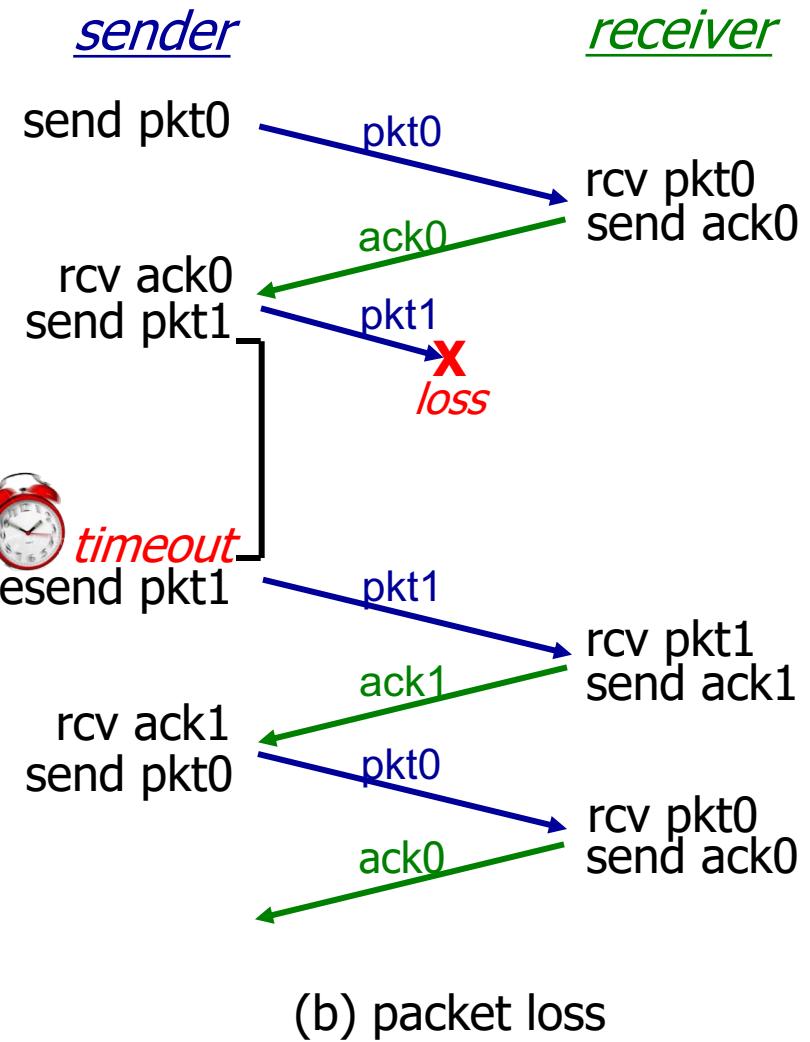
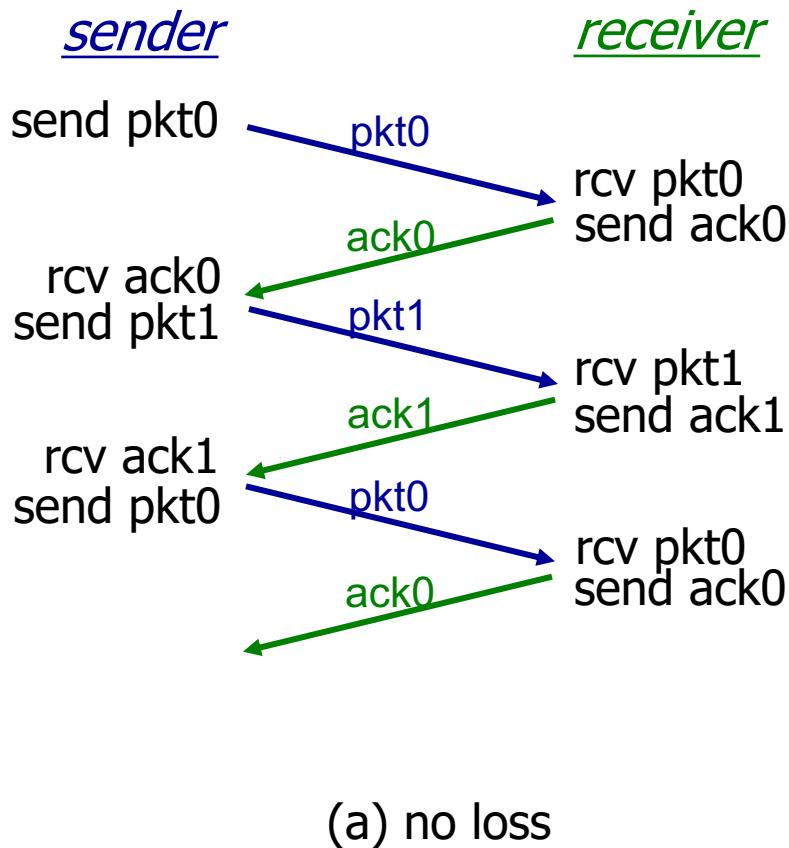
3

2

# rdt3.0 sender

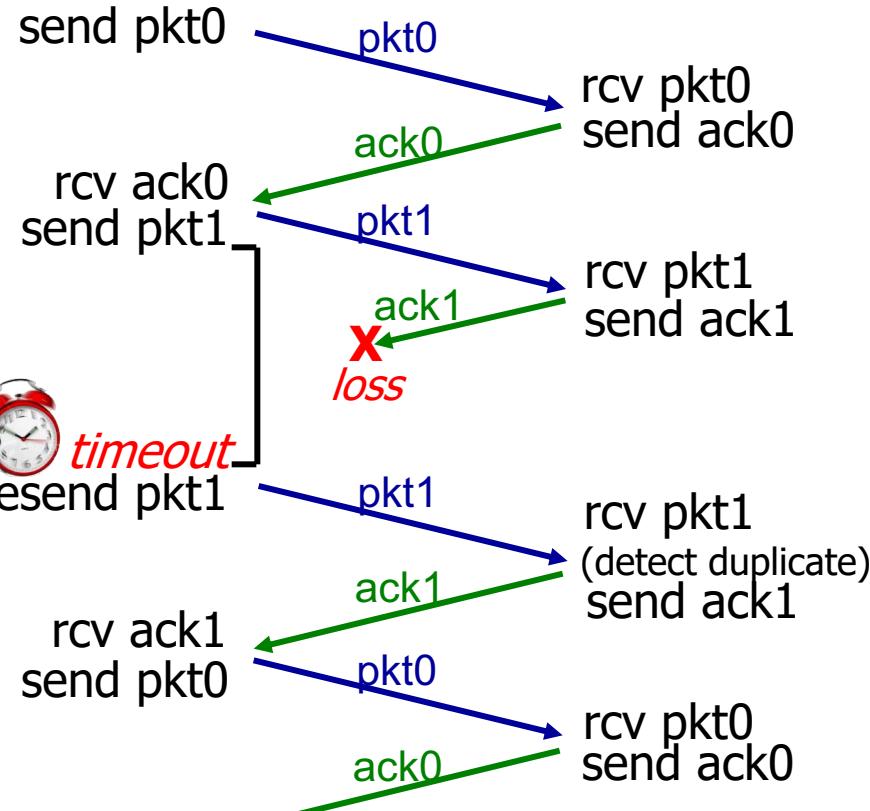


# rdt3.0 in action



# rdt3.0 in action

sender



(c) ACK loss

sender

send pkt0

rcv ack0  
send pkt1

resend pkt1

rcv ack1  
send pkt0

rcv ack0  
send pkt1

pkt0

ack0

pkt1

ack1

pkt1

pkt0

ack1

ack0

pkt1

receiver

rcv pkt0  
send ack0

rcv pkt1  
send ack1

rcv pkt1  
(detect duplicate)  
send ack1

rcv pkt0  
send ack0



**timeout**

rcv ack1

send pkt0

rcv ack1

(do nothing\*\*)

rcv ack0

send pkt1

rcv ack0

send pkt1

rcv ack1

send ack0

rcv ack1

# Performance of rdt3.0

- rdt3.0 is correct, but the performance stinks
- e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

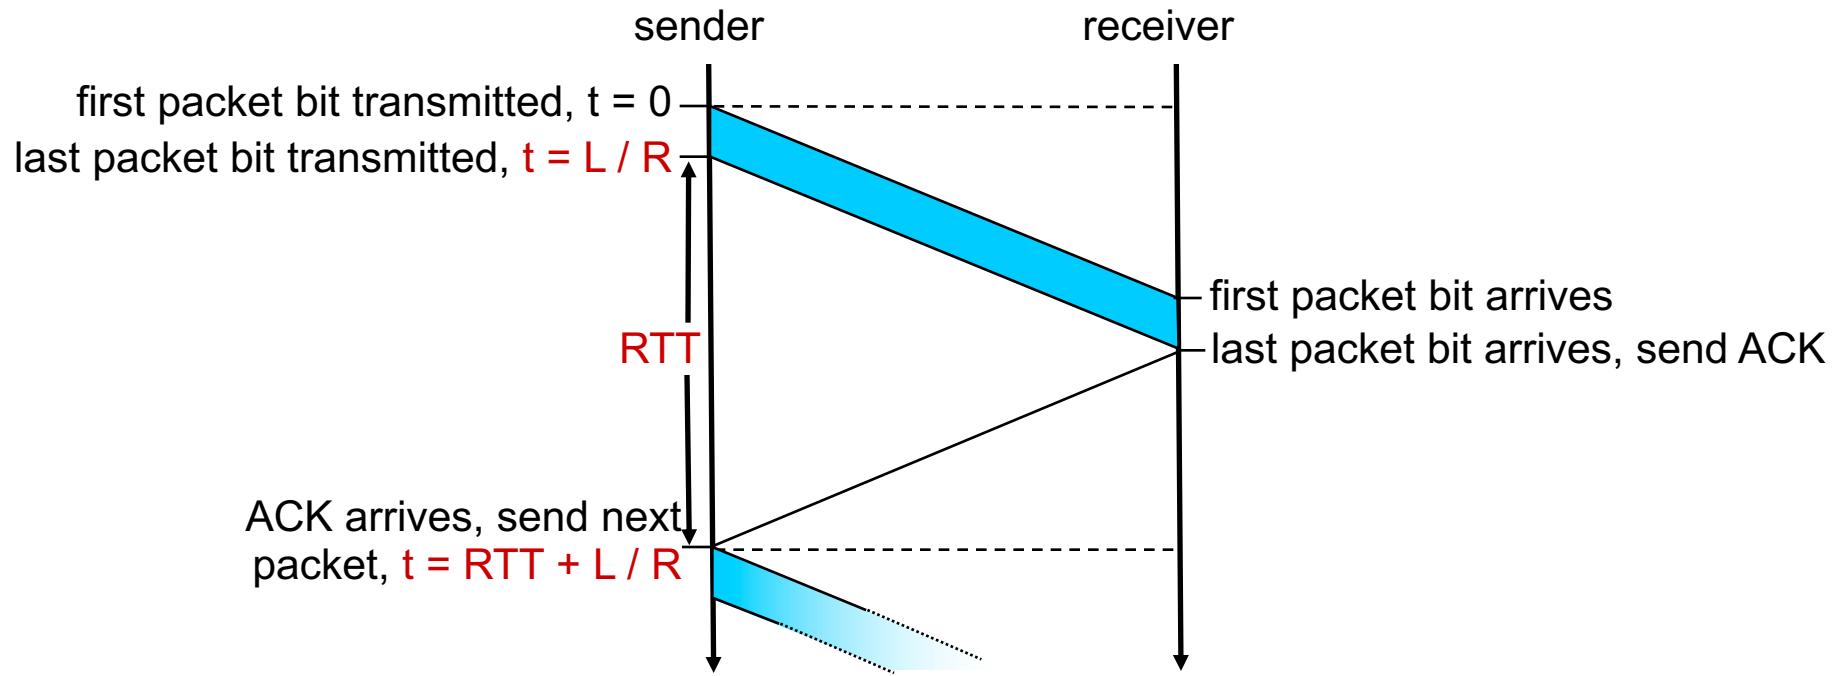
$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

- $U_{\text{sender}}$ : *utilization* – fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

- if RTT=30 msec, 1KB pkt every 30 msec: 33kB/sec thruput over 1 Gbps link
- ❖ **network protocol limits use of physical resources!**

# rdt3.0: stop-and-wait operation



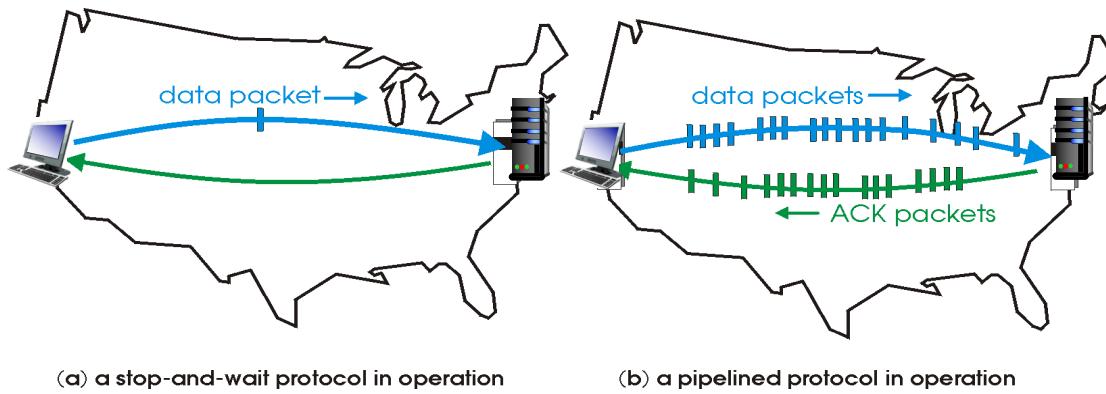
$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

How do we fix this ?

# Pipelined protocols

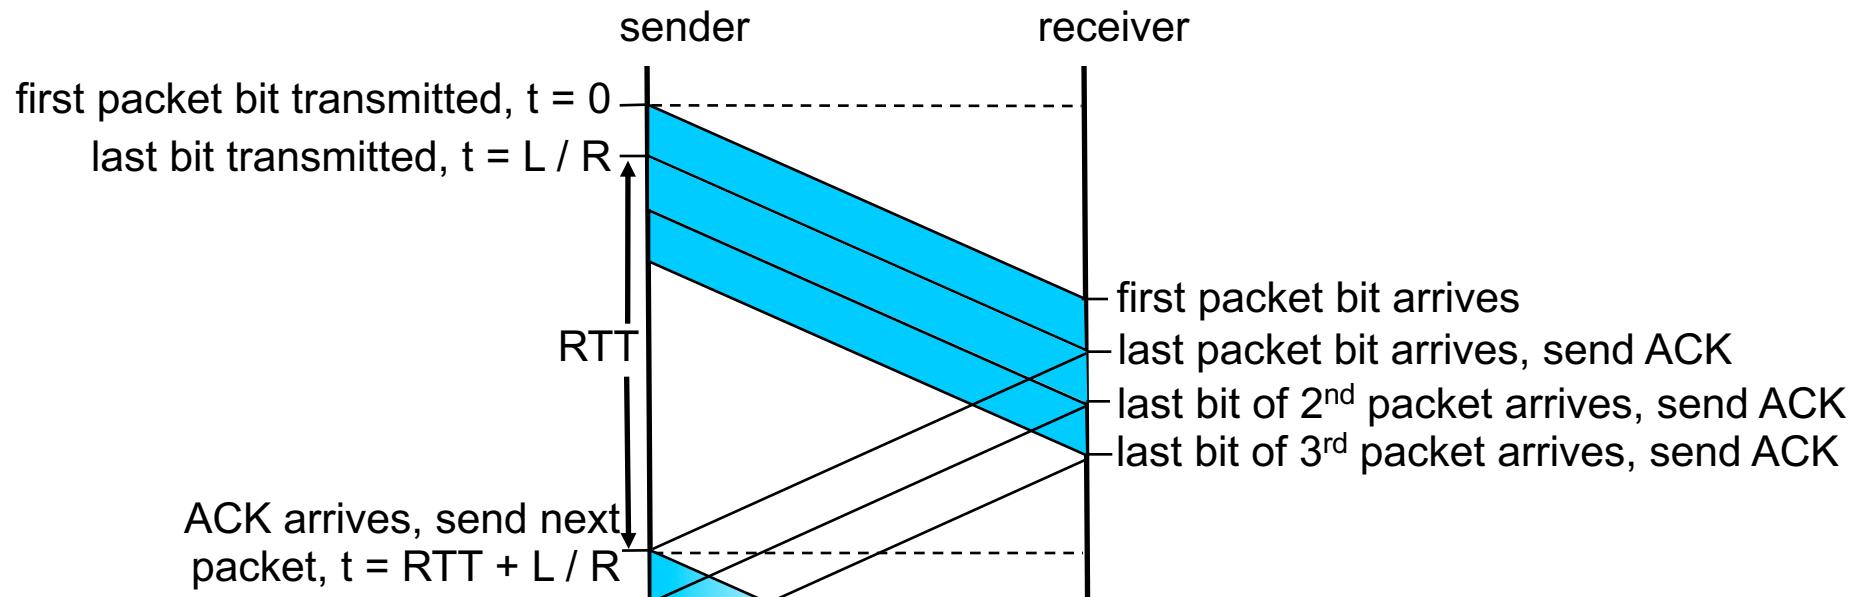
**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver



- two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

# Pipelining: increased utilization



3-packet pipelining increases utilization by a factor of 3!

$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

# Pipelined protocols: overview

## Go-back-N:

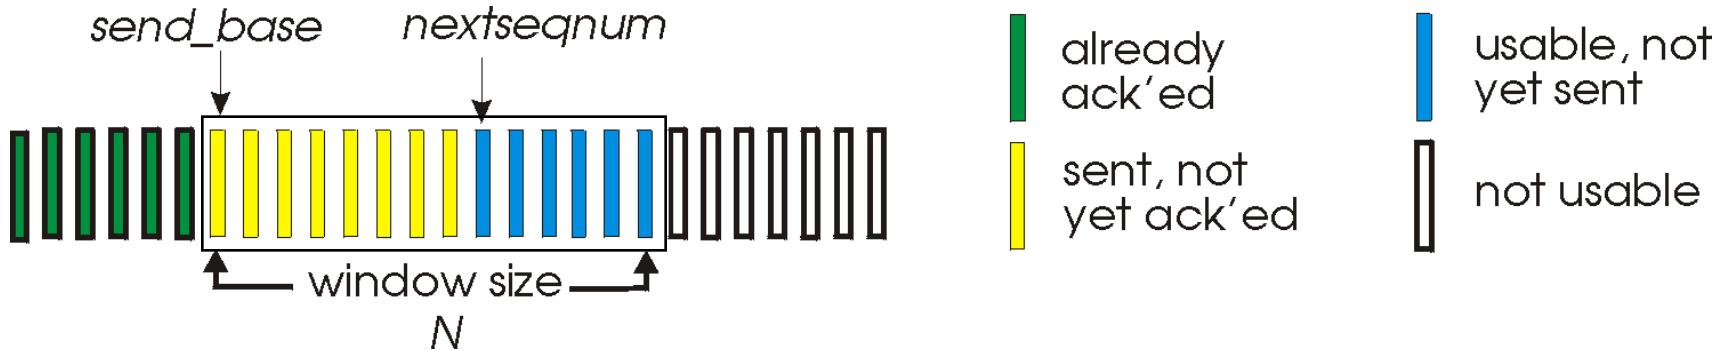
- sender can have up to N unacked packets in pipeline
- receiver only sends *cumulative ack*
  - doesn't ack packet if there's a gap
- sender has timer for oldest unacked packet
  - when timer expires, retransmit *all* unacked packets

## Selective Repeat:

- sender can have up to N unacked'ed packets in pipeline
- rcvr sends *individual ack* for each packet
- sender **maintains timer for each unacked packet**
  - when timer expires, retransmit only that unACKed packet

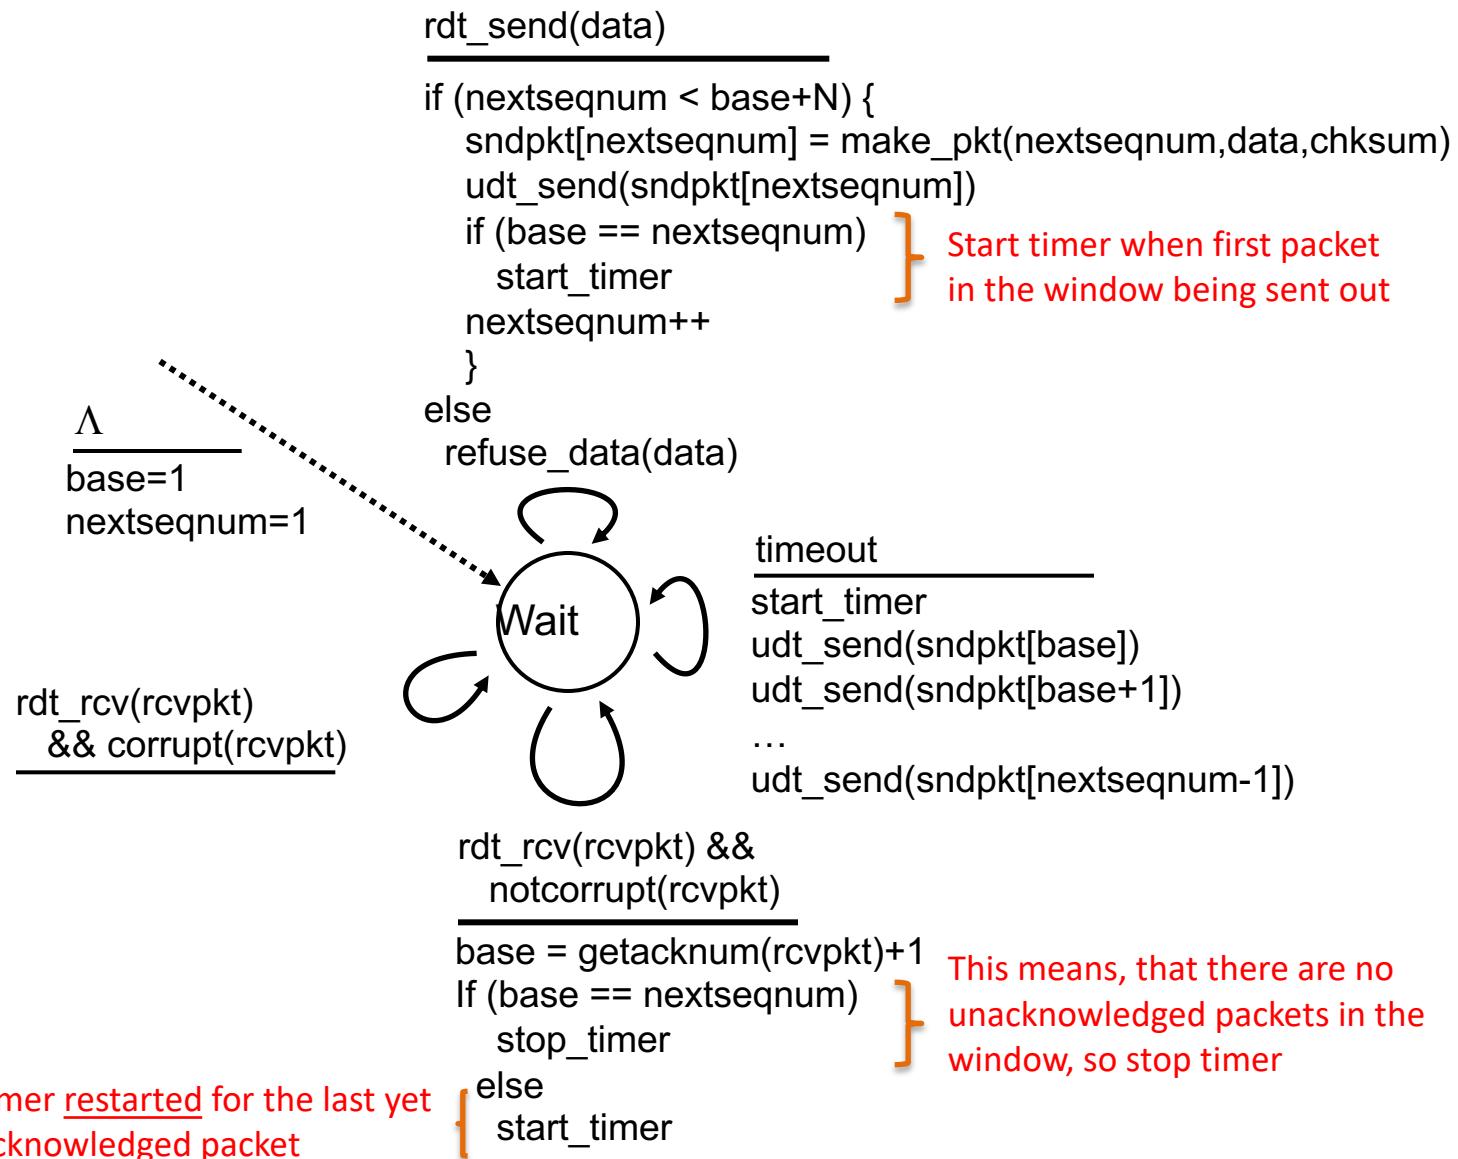
# Go-Back-N: sender

- k-bit seq # in pkt header (modulo  $2^k$  arithmetic for seq #)
- “window” of up to N, consecutive unack’ ed pkts allowed

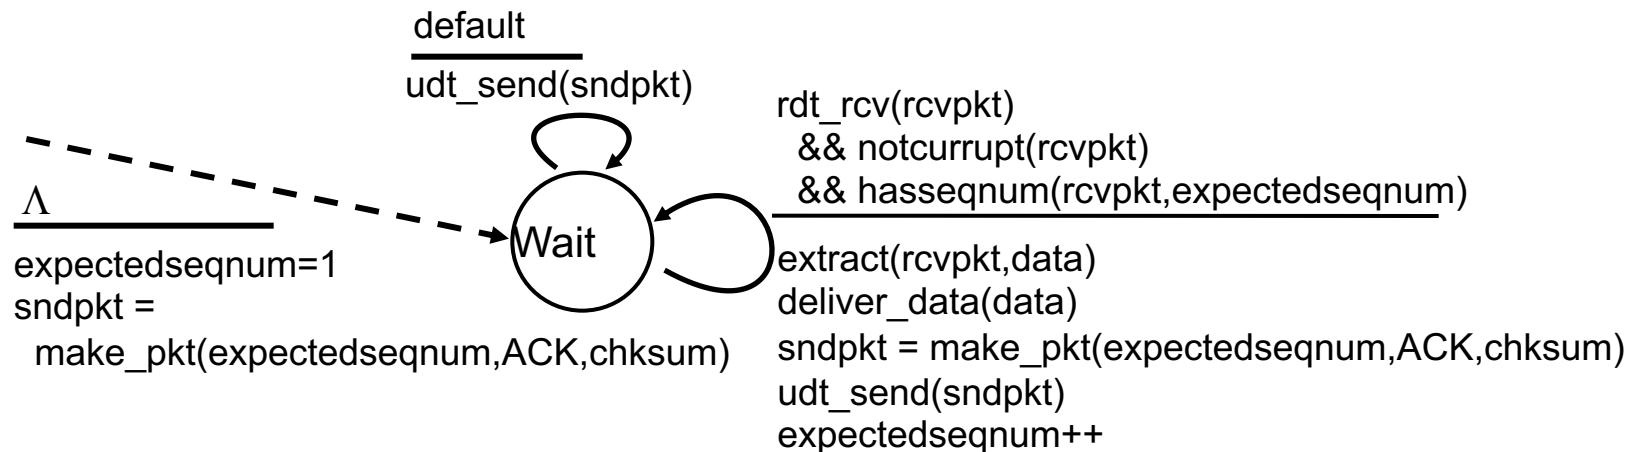


- ❖ ACK(n):ACKs all pkts up to, including seq # n - “*cumulative ACK*”
  - may receive duplicate ACKs (see receiver)
- ❖ timer for **oldest** in-flight pkt
- ❖ *timeout(n)*: retransmit all packets that have been sent but not acknowledged yet

# GBN: sender extended FSM



# GBN: receiver extended FSM



- Maintain variable **expectedseqnum**
- always send ACK if packet received with expected sequence number
- out-of-order pkt:
  - discard (don't buffer): *no receiver buffering! — WHY?*
  - re-ACK last in order seq # packet
  - may generate duplicate ACKs

# GBN in action

sender window (N=4)

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

sender

send pkt0  
send pkt1  
send pkt2  
send pkt3  
(wait)

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

rcv ack0, send pkt4  
rcv ack1, send pkt5

send pkt2  
send pkt3  
send pkt4  
send pkt5

receiver

receive pkt0, send ack0  
receive pkt1, send ack1

Expecting  
seq # 2

receive pkt3, discard,  
(re)send ack1

Expecting  
seq # 2

receive pkt4, discard,  
(re)send ack1

Expecting  
seq # 2

receive pkt5, discard,  
(re)send ack1



*pkt 2 timeout*

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

rcv pkt2, deliver, send ack2  
rcv pkt3, deliver, send ack3  
rcv pkt4, deliver, send ack4  
rcv pkt5, deliver, send ack5

What's bad about this?

